Reviewing material adverse effect clauses in merger agreements

An exploratory study of lightweight analysis and drafting clarity

November 2, 2026

Abstract

This study tests a lightweight analyzer of material adverse effect (MAE) definitions in M&A agreements. It uses 152 public merger-agreement excerpts from MAUD, preserving four expert feature labels per agreement. A custom agreement-disjoint split supports comparisons among literal rules, a majority baseline and a naive Bayes text model. Initial rules achieve 63.7% accuracy on 124 test decisions; the statistical model achieves 74.2%. Corrections informed by test inspection raise rule accuracy to 79.0%, a post hoc result requiring new validation. Errors concern grammatical scope, actor identity and implicit exclusions. The resulting tool supplies exact evidence and drafting questions rather than legal conclusions. The study supports a disciplined review workflow and motivates an independent reviewer experiment; it provides no evidence that suggested edits reduce litigation.

Research question and legal setting

MAE clauses allocate the risk of deterioration between signing and closing. They are a useful narrow setting for studying drafting because their legal effect depends on the relationship among operative triggers, exclusions, exceptions and closing conditions. This pilot asks whether inexpensive, inspectable text analysis can identify selected features and direct attention to specific interpretive choices. It does not classify an agreement as well drafted or predict a lawsuit.

Delaware decisions supply the interpretive lens. Akorn illustrates the importance of the negotiated MAE definition and separate closing conditions [3]. Bardy demonstrates how the chosen comparison group can affect a disproportionate-impact analysis [4]. AB Stable shows why an MAE allocation and an ordinary-course covenant must be read separately and together [5]. These are illustrative decisions, not a systematic case sample. The agreements in MAUD were not screened here for Delaware governing law.

Data and traceability

MAUD is an expert-annotated public merger-agreement dataset published by The Atticus Project under stated CC BY 4.0 terms [1, 2]. The derived sample retains one original MAE excerpt from each of 152 agreements and four Yes/No questions: a forward-looking standard, a closing-ability prong, a stockholder litigation exclusion and a buyer-action exclusion. The resulting 608 labels concern clause features, not defects. Source text, question labels, source row IDs, hashes and transformations remain inspectable.

The audit found no missing labels, duplicate contract-task pairs or empty excerpts. All 152 excerpt texts are distinct after case and whitespace normalization. Five records labeled main by the source contain omission markers; these are preserved and flagged. Full source contracts are obtainable through the pinned downloader, while the evaluated input is an already extracted MAE definition. No OCR or full-contract extraction was tested.

Methods and evaluation

The source CSVs overlap by agreement across different questions. For this four-task combination, a deterministic hash split keeps all questions from an agreement and any normalized exact-text duplicates together: 91 training, 30 development and 31 test agreements. This is a custom pilot, not the published MAUD benchmark. It does not group related drafting templates or stratify labels. The data audit and analysis protocol were developed during implementation; this is not a preregistered study.

The majority baseline uses the most common training answer for each task. The statistical comparator fits one multinomial naive Bayes model per task using alphabetic unigrams and bigrams, token counts capped at three and equal class priors. Smoothing is chosen from 0.1, 1 and 10 using development macro F1 only. The interactive tool uses deterministic matching rules and exact evidence spans. Its separate drafting prompts are not supervised by quality labels.

Evaluation reports each confusion matrix and both class F1 scores, then averages binary macro F1 equally across tasks. Accuracy is calculated over four decisions per agreement. Percentile intervals use 2,000 fixed-seed bootstrap draws of whole test agreements; paired differences resample the same contracts for both methods. This respects within-agreement task dependence but not dependence among drafting families. It does not adjust for test-informed revision.

Method	Accuracy	95% interval	Macro F1
Training majority	64.5%	56.5% to 71.8%	0.389
Initial rules v1	63.7%	55.6% to 71.8%	0.568
Naive Bayes	74.2%	65.3% to 81.5%	0.657
Revised rules v1.1 post hoc	79.0%	72.6% to 85.5%	0.690

Table 1. Custom test sample: 31 agreements and 124 feature decisions. Rule v1.1 was revised after inspecting v1 test errors. The 79.0% result and its interval therefore describe this sample after adaptation; they are not an untouched estimate of generalization. The initial code and results are retained and reproducible.

Revised rule task	TP	TN	FP	FN	Macro F1
Forward looking	15	11	1	4	0.836
Closing ability	12	14	0	5	0.838
Stockholder litigation	7	13	0	11	0.631
Buyer actions	0	26	0	5	0.456

Table 2. Revised-rule counts against expert feature labels. The buyer-action task illustrates class imbalance: 83.9% accuracy comes from predicting no positive cues, with zero recall for the five positive labels. Across tasks, the revised rule advantage over naive Bayes is 4.8 percentage points, with a paired contract-bootstrap interval from -4.0 to 14.5 points. The sample does not establish superiority, even before considering adaptation.

What the errors reveal

Initial rules missed all 17 positive closing-ability labels. Their verb matcher mishandled inflection and spacing and did not cover the common ability-of-the-Company construction. Another rule confused a target action requested by Parent with an action performed by Parent. Correcting these defects reduced total test errors from 45 to 26. The same test data cannot independently validate the improvement.

The revised rules still miss 11 of 18 stockholder-litigation features and all five buyer-action features. An exclusion for claims by any person can encompass shareholder claims without naming shareholders. A closing-ability prong may appear after the first exclusions marker, beyond the rule's assumed boundary. Exact evidence confirms that words occur; it does not establish that their legal role has been understood. Complete error IDs and source excerpts permit further inspection.

Sixteen AI-authored diagnostic fragments add 18 expected feature decisions, explicitly separate from the expert-labeled sample. The revised tool matches 14 expectations and fails four involving negation, incidental co-occurrence or an implicit exclusion. Because the same AI-assisted process produced the probes and expectations, these counts diagnose implementation behavior and provide no independent legal-accuracy estimate.

Drafting principles for further testing

The following principles are case-informed review hypotheses, not causal findings from the dataset. The test runs show why a reviewer must examine grammatical scope and actor identity; they do not prove that any particular wording is defective or should be removed.

1. Separate the operative prongs. Distinguish business deterioration from inability to consummate. State which exclusions and prospective standards apply to each prong, especially when an exception applies only to one branch.

2. Specify the comparison. If an exception restores disproportionate effects, identify the relevant peers, geography, comparison period and approach to measuring excess effects. Bardy illustrates the significance of the selected peer group [4].

3. Choose causal connectors deliberately. Decide whether exclusions reach indirect and mixed causes, then check the use of arising from, resulting from and related to. Consistency is useful only if it preserves the intended allocation.

4. Distinguish a signal from its cause. If missed projections or a share-price decline are excluded, specify whether underlying facts may still count and whether another exclusion controls those facts.

5. Identify the actor. Separate buyer-performed acts, target acts requested by the buyer and target acts merely approved by the buyer. State whose actions and omissions are covered.

6. Coordinate the whole agreement. Check emergency permissions, ordinary-course obligations, consent and notice procedures, conditions and termination rights together. AB Stable cautions against assuming that an MAE exclusion supplies a separate covenant exception [5].

Limits and next study

The evidence supports a transparent research prototype, not an automated legal reviewer. The corpus is small and selected, source omissions persist, party definitions are not resolved, and related templates may cross partitions. Four features cannot summarize an MAE allocation. Common, deliberately negotiated language can trigger every drafting prompt. There are no independent labels for prompt usefulness, ambiguity, enforceability or litigation. Revised-rule intervals quantify sample variation without correcting for test inspection, possible label ambiguity or template dependence.

The next study should freeze a revised implementation before collecting a new agreement-disjoint sample. Two M&A lawyers should independently annotate feature scope and useful drafting issues while blinded to predictions, followed by adjudication. Report feature performance, prompt precision, missed issues, reviewer agreement and review time. Include later filings and identify drafting families. This would test whether the questions help review rather than merely restating common clause topics.

A subsequent randomized reviewer study could compare ordinary review with tool-assisted review, counterbalancing document order and accounting for reviewer and contract dependence. Prespecified outcomes should include lawyer-rated clarity, issue detection, time and unintended changes to commercial allocation. Such a study could estimate effects on review outcomes. A claim about lawsuits would require reliable agreement-to-docket linkage, appropriate timing and censoring, confounder measurement and a credible identification strategy. Clause counts alone cannot identify that effect.

Reproducibility and AI assistance

The repository includes source and derived hashes, labels, split membership, model code, full predictions, error lists and deterministic evaluation commands. Runtime analysis requires Python 3.10 or newer and no external packages or API key. Codex assisted with research, implementation, diagnostics and writing. No LLM inference was used to score the real clauses. Synthetic examples and their expectations are disclosed. Independent legal review and a new frozen test remain outstanding.

References

[1] Wang et al. (2023). MAUD: An Expert-Annotated Legal NLP Dataset for Merger Agreement Understanding. EMNLP.

[2] The Atticus Project. MAUD dataset and CC BY 4.0 license. Accessed November 2 2026.

[3] Akorn, Inc. v. Fresenius Kabi AG, 2018 WL 4719347 (Del. Ch. Oct. 1, 2018). See opinion pp. 123-125 on the MAE condition and definition.

[4] Bardy Diagnostics, Inc. v. Hill-Rom, Inc., 2021 WL 2886188 (Del. Ch. July 9, 2021). See opinion pp. 89-93 on the comparison group.

[5] AB Stable VIII LLC v. MAPS Hotels and Resorts One LLC, 268 A.3d 198 (Del. 2021). See the discussion of MAE risk allocation and the ordinary-course covenant.